
Backbone.js on Rails

Table of Contents

1. Preface (section unstated)	2
2. Getting up to speed (section unstated)	2
2.1. Backbone.js online resources	2
2.2. JavaScript online resources and books	2
3. Introduction (section unstated)	2
3.1. Why use Backbone.js	2
3.2. When not to use Backbone.js	2
3.3. Why not SproutCore, Cappuccino, Knockout.js, Spine, etc.	2
3.4. The Example Application	2
4. Organization	3
4.1. Backbone.js and MVC	3
4.2. What Goes Where	4
4.3. Namespacing your application	5
4.4. Mixins	5
5. Rails Integration	6
5.1. Organizing your Backbone.js code in a Rails app	6
5.2. Rails 3.0 and prior	6
5.3. Rails 3.1	8
5.4. An Overview of the Stack: Connecting Rails and Backbone.js	10
5.5. Customizing your Rails-generated JSON	15
5.6. Converting an existing page/view area to use Backbone.js	17
5.7. Automatically using the Rails authentication token	22
6. Routers, Views, and Templates	24
6.1. View explanation (chapter unstated)	24
6.2. Routers (chapter unstated)	24
6.3. Templating strategy (chapter unstated)	24
6.4. View helpers (chapter unstated)	24
6.5. Form helpers (chapter unstated)	24
6.6. Event binding	24
6.7. Cleaning Up: Unbinding	26
6.8. Swapping router	32
6.9. Composite views	34
6.10. How to use multiple views on the same model/collection (chapter unstated)	39
6.11. Internationalization (chapter unstated)	39
7. Models and collections	39
7.1. Naming conventions (chapter unstated)	39
7.2. Nested resources (chapter unstated)	39
7.3. Model associations	39
7.4. Filters and sorting	41
7.5. Client/Server duplicated business logic (chapter unstated)	45
7.6. Validations (chapter unstated)	45
7.7. Synchronizing between clients (chapter unstated)	45
8. Testing (section unstated)	45
8.1. Full-stack integration testing	45

8.2. Isolated unit testing	45
9. The JavaScript language (section unstated)	45
9.1. Model attribute types and serialization	45
9.2. Context binding (JS <code>this</code>)	45
9.3. CoffeeScript with Backbone.js	45
10. Security (stub)	45
10.1. XSS with JSON bootstrapping (stub)	45
10.2. XSS in HTML templates (stub)	46
11. Performance (stub)	46
11.1. Dependency choice	46

1. Preface (section unstated)

2. Getting up to speed (section unstated)

2.1. Backbone.js online resources

2.2. JavaScript online resources and books

3. Introduction (section unstated)

3.1. Why use Backbone.js

3.2. When not to use Backbone.js

3.3. Why not SproutCore, Cappuccino, Knockout.js, Spine, etc.

3.4. The Example Application

Rails 3.1.0.rc5

Ruby 1.9.2

Backbone.js and Underscore.js are the non-minified versions. This is for informational purposes, but also because the Rails 3.1 asset pipeline will compress and minify them.


While Rails 3.1 defaults to CoffeeScript, we have decided to make all example code normal Javascript as we believe that will be the most understandable to the current readers.

4. Organization

4.1. Backbone.js and MVC

Model–View–Controller (MVC) is an architectural pattern used in many applications to isolate "domain logic" (the application logic for the user) from the user interface (input and presentation).

Figure 1. Model-view-controller concept



In the above diagram a solid line represents a direct association and a dashed line represents an indirect association (for example, via an observer).

As a user of Rails, you're likely already familiar with the concept of MVC and the benefits that the separation of concerns can give you. However, Rails itself is not doing "traditional" MVC. A traditional MVC is event-based. This means that the views trigger events which the controller figures out what to do with. It can be argued that the requests generated by the browser are the "events" in Rails, however, due to the single-threaded, request-response nature of the web, the control flow between the different levels of MVC is much more straightforward.

Given that Javascript has events, and that much of the interactions between the different components of Backbone.js in the browser are not limited to request/response, programming with Backbone.js is in a lot of ways more like working with a traditional MVC architecture.

That being said, technically speaking, Backbone.js is *not* MVC, and this was acknowledged by the creators of Backbone.js when they renamed Controllers to Routers in version 0.5.0.

So what is Backbone.js then if not MVC? Technically speaking, it's just the Models and the Views with a Router to handle flow between them. In Backbone.js the views will handle many of the aspects that controllers would typically handle, such as actually figuring out what to do next and what to render.

While you could do it, the benefit of actually introducing a Controller in your application would be limited, and the more pragmatic approach is to realize the great organization that Backbone.js

gives you is much better than what you had before. The fact that it doesn't have a nice name, or strict adherence to a pattern, isn't worth worrying about.

4.2. What Goes Where

Part of the initial learning curve of Backbone.js can be figuring out what goes where, and mapping it to your expectations set by working with Rails. In Rails we have Models, Views, Controllers, and Routers. In Backbone.js, we have Models, Collections, Views, Templates, and Routers.

The models in Backbone.js and Rails are analogous. Backbone.js collections are just ordered sets of models. Because it lacks controllers, Backbone.js routers and views work together to pick up the functionality provided by Rails controllers. Finally, in Rails, when we say views, we actually mean templates. In Backbone.js, however, you have a separation between the view and templates.

Once you introduce Backbone.js into your stack, you grow the layers in your stack by four levels. This can be daunting at first, and frankly, at times it can be difficult to keep everything going on in your application straight. Ultimately, the additional organization and functionality of Backbone.js outweighs the costs, so let's break it down.

Rails

- Model
- Controller
- View

Backbone.js

- Model and Collection
- Router
- View
- Template

In a typical Rails and Backbone.js application, the initial interaction between the layers will be as follows:

- A request from a user comes in the **Rails router** identifies what should handle the request based on the URL
- The **Rails controller action** to handle the request is called, some initial processing may be performed
- The **Rails view template** is rendered and returned to the user's browser
- The **Rails view template** will include **Backbone.js initialization**, usually this is populating some **Backbone collections** as sets of **Backbone models** with JSON data provided by the **Rails view**
- The **Backbone.js router** determines which of its methods should handle the display based on the URL

- The **Backbone.js router** method calls that method, some initial processing may be performed, and one or more **Backbone.js views** are rendered
- The **Backbone.js view** reads **templates** and uses **Backbone.js** models to render itself onto the page

At this point, the user will see a nice page in their browser and be able to interact with it. The user interacting with elements on the page will trigger actions to be taken at any level of the above sequence: **Backbone.js model**, **Backbone.js views**, **Backbone.js router**, or requests to the remote server.

Requests to the remote server may be any one of the following:

- At the **Backbone.js model** or **Backbone.js collection** level, communicating with Rails via JSON.
- Normal Ajax requests, not using Backbone.js at all.
- Normal requests that don't hit Backbone.js and trigger a full page reload.

Which of the above remote server interactions you use will depend upon the desired result, and the type of user interface. This book should help you understand which interaction you'll want to choose for each portion of your application.

4.3. Namespacing your application

You will want to create an object in Javascript for your Backbone.js application to reside. This variable will serve as a namespace for your Backbone.js application. Namespacing all of the Javascript is desirable to avoid potential collisions in naming. For example, it's possible that a Javascript library you want to use might also create a Task variable. If you didn't namespace your Task model then this would conflict.

This variable includes a place to hold Models, Collections, Views, and Routes, and an init method which will be called to initialize the application. Its very common to create a new Router in the init function, and Backbone.history.start() must be called in order to route the initial URL. This app variable will look like the following.

```
var ExampleApp = {
  Models: {},
  Collections: {},
  Views: {},
  Routers: {},
  init: function() {
    new ExampleApp.Routers.Tasks();
    Backbone.history.start();
  }
};
```

You can find this file in the example app in `app/assets/javascripts/example_app.js`.

4.4. Mixins

Backbone provides a basic mechanism for inheritance. Often you'll want to build a collection of related, reusable behavior and include that in several classes that already inherit from a Backbone base class. In these cases, you'll want to use a mixin [<http://en.wikipedia.org/wiki/Mixin>].

Backbone includes `Backbone.Events` [<http://documentcloud.github.com/backbone/#Events>] as an example of a mixin.

Here, we create a mixin named `Observer` that contains behavior for binding to events in a fashion that can be cleaned up later:

```
var Observer = {
  bindTo: function(source, event, callback) {
    source.bind(event, callback, this);
    this.bindings.push({ source: source, event: event, callback: callback });
  },

  unbindFromAll: function() {
    _.each(this.bindings, function(binding) {
      binding.source.unbind(binding.event, binding.callback);
    });
    this.bindings = [];
  }
};
```

We can mix `Observer` into a class by using Underscore's `_.extend` on the prototype of that class:

```
SomeCollectionView = Backbone.Collection.extend({
  initialize: function() {
    this.bindTo(this.collection, "change", this.render);
  },

  leave: function() {
    this.unbindFromAll(); // calling a method defined in the mixin
    this.remove();
  }
});

_.extend(SomeCollectionView.prototype, Observer);
```

5. Rails Integration

5.1. Organizing your Backbone.js code in a Rails app

When using Backbone.js in a Rails app, you'll have two kinds of Backbone.js-related assets: classes and templates.

5.2. Rails 3.0 and prior

With Rails 3.0 and prior, store your Backbone.js classes in `public/javascripts`:

```
public/
  javascripts/
    jquery.js
    jquery-ui.js
  collections/
    users.js
    todos.js
  models/
```

```
user.js
todo.js
routers/
  users_router.js
  todos_router.js
views/
  users/
    users_index.js
    users_new.js
    users_edit.js
  todos/
    todos_index.js
```

If you are using templates, we prefer storing them in `app/templates` to keep them separated from the server views:

```
app/
  views/
    pages/
      home.html.erb
      terms.html.erb
      privacy.html.erb
      about.html.erb
  templates/
    users/
      index.jst
      new.jst
      edit.jst
    todos/
      index.jst
      show.jst
```

On Rails 3.0 and prior apps, we use Jammit for packaging assets and precompiling templates:

<http://documentcloud.github.com/jammit/>
<http://documentcloud.github.com/jammit/#jst>

Jammit will make your templates available in a top-level JST object. For example, to access the above `todos/index.jst` template, you would refer to it as:

```
JST['todos/index']
```

Variables can be passed to the templates by passing a Hash to the template, as shown below.

```
JST['todos/index']({ model: this.model })
```

Note

A note on Jammit and a JST naming gotcha

One issue with Jammit that we've encountered and worked around is that the JST template path can change when adding new templates.

When using Jammit, there is a slightly sticky issue as an app grows from one template subdirectory to multiple template subdirectories.

Let's say you place templates in `app/templates`. You work for a while on the "Tasks" feature, placing templates under `app/templates/tasks`. So, `window.JST` looks something like:

```
JST[ 'form' ]  
JST[ 'show' ]  
JST[ 'index' ]
```

Now, you add another directory under `app/templates`, say `app/templates/user`. Now, all JST references are prefixed with their parent directory name so they are unambiguous:

```
JST[ 'tasks/form' ]  
JST[ 'tasks/show' ]  
JST[ 'tasks/index' ]  
JST[ 'users/new' ]  
JST[ 'users/show' ]  
JST[ 'users/index' ]
```

This breaks existing JST references. You can work around this issue by applying the following monkeypatch to Jammit, in `config/initializers/jammit.rb`

```
Jammit::Compressor.class_eval do  
  private  
  def find_base_path(path)  
    File.expand_path(Rails.root.join('app', 'templates'))  
  end  
end
```

As applications are moving to Rails 3.1, they're also moving to Sprockets for the asset packager. Until then, many apps are using Jammit for asset packaging. We have an open issue and workaround:

<https://github.com/documentcloud/jammit/issues/192>

5.3. Rails 3.1

Rails 3.1 introduces the asset pipeline:

http://edgeguides.rubyonrails.org/asset_pipeline.html

which uses the Sprockets library for preprocessing and packaging assets:

<http://getsprockets.org/>

To take advantage of the built-in asset pipeline, organize your Backbone.js templates and classes in paths available to the asset pipeline. Classes go in `app/assets/javascripts/`, and templates go alongside, in `app/assets/templates/`:

```
app/  
  assets/  
    javascripts/  
      collections/  
        todos.js  
    models/  
      todo.js  
    routers/  
      todos_router.js  
    views/  
      todos/  
        todos_index.js
```



```
templates/  
  todos/  
    index.jst.ejs  
    show.jst.ejs
```

In Rails 3.1, jQuery is provided by the `jquery-rails` gem, and no longer needs to be included in your directory structure.

Using Sprockets' preprocessors, we can use templates as before. Here, we're using the EJS template preprocessor to provide the same functionality as Underscore.js' templates. It compiles the `*.jst` files and makes them available on the client side via the `window.JST` object. Identifying the `.ejs` extension and invoking EJS to compile the templates is managed by Sprockets, and requires the `ejs` gem to be included in the application Gemfile.

Note

Underscore.js templates: <http://documentcloud.github.com/underscore/#template>

EJS gem: <https://github.com/sstephenson/ruby-ejs>

Sprockets support for EJS: https://github.com/sstephenson/sprockets/blob/master/lib/sprockets/ejs_template.rb

To make the `*.jst` files available and create the `window.JST` object, require them in your `application.js` Sprockets manifest:

```
// other application requires  
//= require_tree ../templates  
//= require_tree .
```

Additionally, load order for Backbone.js and your Backbone.js app is very important. jQuery and Underscore.js must be loaded before Backbone.js, then the Rails authenticity token patch must be applied. Then your models must be loaded before your collections (because your collections will reference your models) and then your routers and views must be loaded.

Fortunately, sprockets can handle this load order for us. When all is said and done your `application.js` Sprockets manifest will be as shown below.

```
//= require jquery  
//= require jquery_ujs  
//  
//= require underscore  
//= require backbone  
//= require backbone.authtokenadapter  
//  
//= require example_app  
//  
//= require_tree ./models  
//= require_tree ./collections  
//= require_tree ./views  
//= require_tree ./routers  
//= require_tree ../templates  
//= require_tree .
```

The above is taken from the example application included with this book. You can view it at `example_app/app/assets/javascripts/application.js`.

5.4. An Overview of the Stack: Connecting Rails and Backbone.js

By default Backbone.js communicates with your Rails application via JSON gets and posts. If you've ever made a JSON API for your Rails app, then for the most part this will be very similar.

If you haven't ever made a JSON API for your Rails application before, lucky you, it's pretty straightforward.

5.4.1. Setting Up Rails Models

One important aspect to keep in mind as you plan out how your Backbone.js interface will behave, and how it will use your Rails back-end is that there is no need to have a one-to-one mapping between your Rails models and your Backbone.js models.

The smaller an application is, the more likely that there will be a one-to-one mapping between both Backbone.js and Rails models and controllers.

However, if you have a sufficiently complex application, it's more likely that you *won't* have a one-to-one mapping due to the differences in the tools Backbone.js gives you and the fact that you're building a user-interface, not a back-end. Some of the reasons why you won't have a one to one mapping include:

- Because you're building a user interface, not a back-end, it's likely that some of your backbone models will aggregate information from multiple Rails models into one Backbone.js model.
- This Backbone.js model may or may not be named the same as one of your Rails models.
- Backbone.js gives you a new type of object not present in Rails: Collections.
- Backbone.js doesn't have the concept of relationships out of the box.

With that said, let's take the simple case first and look at how you might make a Backbone.js version of a Rails model.

In our example application, we have a Task model. The simplest Backbone.js representation of this model would be as shown below.

```
var Task = Backbone.Model.extend({  
  urlRoot: '/tasks'  
});
```

The urlRoot property above indicates to Backbone.js that the server url for instances of this model will be found at /tasks/:id.

In Rails, it's possible to access individual Tasks, as well as all Tasks (and query all tasks) through the same Task model. However, in Backbone.js models only represent the singular representation of a Task. Backbone.js splits out the plural representation of Tasks into what it calls Collections.

The simplest Backbone.js collection to represent our Tasks would be the following.

```
var Tasks = Backbone.Collection.extend({  
  model: Task  
});
```

If we specify the url for Tasks in our collection instead, then models within the collection will use the collection's url to construct their own urls, and the urlRoot no longer needs to be specified in the model. If we make that change, then our collection and models will be as follows.

```
var Tasks = Backbone.Collection.extend({
  model: Task,
  url: '/tasks'
});

var Task = Backbone.Model.extend({});
```

Notice in the above model definitions that there is no specification of the attributes on the model. Like ActiveRecord, Backbone.js models get their attributes from the schema and data given to them. In the case of Backbone.js, this schema and data are the JSON from the server.

The default JSON representation of an ActiveRecord model is a Hash that includes all the model's attributes. It does not include the data for any related models or any methods on the model, but it does include the ids of any related models as those are stored in a `relation_name_id` attribute on the model.

The JSON representation of your ActiveRecord models will be retrieved by calling `to_json` on them. You customize the output of `to_json` by overriding the `as_json` method in your model. We'll touch on this more later in the section "Customizing your Rails-generated JSON."

5.4.2. Setting Up Rails Controllers

The Backbone models and collections will talk to your Rails controllers. While your models may not have a one-to-one mapping with their Rails counterparts, it is likely that you'll have at least one controller corresponding to every Backbone.js model.

Fortunately for us, Backbone.js models will communicate in the normal RESTful way that Rails controllers understand, using the proper verbs to support the standard RESTful Rails controller actions: index, show, create, update, and destroy. Backbone.js does not make any use the new action.

Therefore, it's just up to us to write a *normal* restful controller.

There are a few different ways you can write your controllers for interacting with you Backbone.js models and collections. However, the newest and cleanest way is to use the `respond_with` method introduced in Rails 3.0.

When using `respond_with`, in your controller you specify what formats are supported with the method `respond_to`. In your individual actions, you then specify the resource or resources to be delivered using `respond_with`, as shown in the example Tasks controller and index action below.

```
class TasksController < ApplicationController::Base
  respond_to :html, :json

  def index
    respond_with(@tasks = Task.all)
  end
end
```

In the above example Tasks controller, the `respond_to` line declares that this controller should respond to both the HTML and JSON formats. Then, in the index action, the `respond_with` call will perform the appropriate action for the requested format.

The above controller is equivalent to the following one, using the older `respond_to` method.

```
class TasksController < ApplicationController::Base
  def index
    @tasks = Task.all
    respond_to do |format|
      format.html
      format.json { render :json => @tasks }
    end
  end
end
```

Using `respond_with` you can create succinct controllers that respond with a normal web page, but also expose a JSON api that Backbone.js will use.

5.4.2.1. Validations and your HTTP API

If a Backbone.js model has a `validate` method defined, it will be validated before its attributes are set. If validation fails, no changes to the model will occur, and the "error" event will be fired. Your `validate` method will be passed the attributes that are about to be updated. You can signal that validation passed by returning nothing from your `validate` method. You can signify that validation has failed by returning something from the method. What you return can be as simple as a string, or a more complex object that describes the error in all its gory detail.

In practice, much of the validation logic for your models will continue to be handled on the server, as fully implementing validations on the client side would often require duplicating a lot of server-side business logic.

TODO: Is it possible to smoothly integrate Backbone.js and the `client_side_validations` gem?

Instead, your Backbone.js applications will likely rely on server-side validation logic. How to handle a failure scenario is passed in to Backbone.js model `save` call as a callback, as shown below.

```
task.save({title: "New Task title"}, {
  error: function(){
    // handle error from server
  }
});
```

The error callback will be triggered if your server returns a non-200 response. Therefore, you'll want your controller to return a non-200 HTTP response code if validations fail.

A controller that does this would be as shown in the following example.

```
class TasksController < ApplicationController::Base
  respond_to :json

  def create
    @task = Task.new(params[:task])
    if @task.save
      respond_with(@task)
    else
      respond_with(@task, :status => :unprocessable_entity)
    end
  end
end
```

Your error callback will receive both the model as it was attempted to be saved and the response from the server. You can take that response and handle the errors returned by the above controller in whatever way is fit for your application. For more information about handling and displaying errors, see the Form helpers section of the Views and Templates chapter.

5.4.3. Setting Up Views

Most Backbone.js applications will be a "single-page app". This means that your Rails application will render a single-page which properly sets up Backbone.js and the data it will use. From there, ongoing interaction with your Rails application occurs via the JSON apis.

The most common page for this single-page application will be the index action of a controller, as in our example application and the tasks controller.

You will want to create an object in Javascript for your Backbone.js application to reside. For more information on this namespacing see the "Namespacing your application" section of the Organization chapter.

This namespace variable holds your Backbone.js application's Models, Collections, Views, and Routes, and has an init method which will be called to initialize the application.

This namespace variable will look like the following.

```
var ExampleApp = {
  Models: {},
  Collections: {},
  Views: {},
  Routers: {},
  init: function() {
    new ExampleApp.Routers.Tasks();
    Backbone.history.start();
  }
};
```

You can find this file in the example app in `app/assets/javascripts/example_app.js`.

Important

You must instantiate a Backbone.js router before calling `Backbone.history.start()` otherwise `Backbone.history` will be undefined.

Then, inside `app/views/tasks/index.html.erb` you will call the initialize method. This will appear as follows.

```
<%= content_for :javascript do -%>
  <%= javascript_tag do %>
    ExampleApp.init();
  <% end %>
<% end -%>
```

For performance reasons, you will almost always "prime the pump" and give Backbone.js its initial data within the HTML view for this page. In our example, the tasks have already been provided to the view in a `@tasks` instance variable, and that can be used to prime the pump, as shown below.

```
<%= content_for :javascript do -%>
```

```
<%= javascript_tag do %>
  ExampleApp.init(<%= @tasks.to_json %>);
<% end %>
<% end -%>
```

The above example uses Erb to pass the JSON for the tasks to the init method.

Once you make this change, the ExampleApp.init method then becomes:

```
var ExampleApp = {
  Models: {},
  Collections: {},
  Views: {},
  Routers: {},
  init: function(tasks) {
    new ExampleApp.Routers.Tasks();
    this.tasks = new ExampleApp.Collections.Tasks(tasks);
    Backbone.history.start();
  }
};
```

Finally, you must have a Router in place which knows what to do. We'll cover routers in more detail in the Views and Templates chapter. For a more in-depth presentation on writing and using routes please go there. However, routers are an important part of the infrastructure you need to start using Backbone.js and we can't make our example here work without them.

Backbone.js routers provide methods for routing application flow based on client-side URL fragments (#fragment).

```
ExampleApp.Routers.Tasks = Backbone.Router.extend({
  routes: {
    "": "index"
  },

  index: function() {
    // We've reached the end of Rails integration - it's all Backbone from here!

    alert('Hello, world! This is a Backbone.js router action.');
```

// Normally you would continue down the stack, instantiating a
// Backbone.View class, calling render() on it, and inserting its element
// into the DOM.

```
  }
});
```

A basic router consists of a routes hash which is a mapping between url fragments and methods on the router. If the current URL fragment, or one that is being visited matches one of the routes in the hash, its method will be called.

The example router above is all that is needed to complete our Backbone.js infrastructure. When a user visits /tasks the index.html.erb view will be rendered which properly initialized Backbone.js and its dependencies and the Backbone.js models, collections, routers, and views.

Note

Backbone.js now includes support for pushState, which can use real, full URLs instead of url fragments for routing.

However, `pushState` support in Backbone.js is fully opt-in due to lack of browser support and that additional server-side work is required to support it.

`pushState` support is current limited to the latest versions of Firefox, Chrome, and Safari and Mobile Safari. For a full listing of support and more information about the History API, of which `pushState` is a part, visit <http://diveintohtml5.org/history.html#how>

Thankfully, if you opt-in to `pushState` in Backbone.js, browsers that don't support `pushState` will continue to use hash-based URL fragments, and if a hash URL is visited by a `pushState`-capable browser, it will be transparently upgraded to the true URL.

In addition to browser support, another hurdle to seamless use of `pushState` is that because the URL used are real URLs, your server must now how to render each of the urls. For example, if your Backbone.js application has a route of `/tasks/1`, your server-side application must be able to respond to that page if the browser visits that URL directly.

For most applications, you can handle this by just rendering the content you would have for the root URL and letting Backbone.js handle the rest of the routing to the proper location. But for full search-engine crawlability, your server-side application will need to render the entire HTML of the requested page.

5.5. Customizing your Rails-generated JSON

There are a few common things you'll do in your Rails app when working with Backbone.js.

First, it's likely that you'll want to switch from including all attributes (the default) to delivering some subset.

This can be done by specifying explicitly only the attributes that are to be included (whitelisting), or specifying the attributes that should *not* be included (blacklisting). Which one you choose will depend on how many attributes your model has and how paranoid you are about something important appearing in the JSON when it shouldn't be there.

If you're concerned about sensitive data unintentionally being included in the JSON when it shouldn't be then you'll want to whitelist, to switch to everything being explicitly included in the JSON with the `:only` option:

```
def as_json(options = {})
  super(options.merge(:only => [ :id, :title ]))
end
```

The above `as_json` override will make it so that the JSON will *only* include the `id` and `title` attributes, even if there are many other attributes on the model.

If instead you want to include all attributes by default and just exclude a few, you accomplish this with the `:except` option:

```
def as_json(options = {})
  super(options.merge(:except => [ :encrypted_password ]))
end
```

Another common customization you will want to do in the JSON is include the output of methods (say, calculated values) on your model. This is accomplished with the `:methods` option, as shown in the following example.

```
def as_json(options = {})
  super(options.merge(:methods => [ :calculated_value ]))
end
```

The final thing you'll most commonly do with your JSON is include related objects. If the `Task` model has `many :comments`, include all of the JSON for comments in the JSON for a `Task` with the `:include` option:

```
def as_json(options = {})
  super(options.merge(:include => [ :comments ]))
end
```

As you probably suspect, you can then customize the JSON for the comments by overriding the `as_json` method on the `Comment` model.

While these are the most common `as_json` options you'll use when working with Backbone.js, it certainly isn't all of them. The official, complete, documentation for the `as_json` method can be found here: http://apidock.com/rails/ActiveModel/Serializers/JSON/as_json

5.5.1. ActiveRecord::Base.include_root_in_json

Depending on the versions, Backbone.js and Rails may have different expectations about the format of JSON structures; specifically, whether or not a root key is present. When generating JSON from Rails, this is controlled by the `ActiveRecord` setting `ActiveRecord::Base.include_root_in_json`.

```
> ActiveRecord::Base.include_root_in_json = false
> Task.last.as_json
=> {"id"=>4, "title"=>"Enjoy a three mile swim"}

> ActiveRecord::Base.include_root_in_json = true
> Task.last.as_json
=> {"task"=>{"id"=>4, "title"=>"Enjoy a three mile swim"}}
```

In Rails 3.0, `ActiveRecord::Base.include_root_in_json` is set to `true`. In 3.1, it defaults to `false`. This reversal was made to simplify the JSON returned by default in Rails application, but it is fairly big change from the default behavior of Rails 3.0.

Practically speaking, this change is a good one, but take particular note if you're upgrading an existing Rails 3.0 application to Rails 3.1 and you already have a published API; you may need to expose a new version of your API.

From the Backbone.js side, the default behavior expects no root node. This behavior is defined in a few places: `Backbone.Collection.prototype.parse`, `Backbone.Model.prototype.parse`, and `Backbone.Model.prototype.toJSON`:

```
_.extend(Backbone.Collection.prototype, Backbone.Events, {
  // http://documentcloud.github.com/backbone/#Collection-parse
  parse : function(resp, xhr) {
    return resp;
  },

  // snip...
});

_.extend(Backbone.Model.prototype, Backbone.Events, {
```



```
// http://documentcloud.github.com/backbone/#Model-toJSON
toJSON : function() {
  return _.clone(this.attributes);
},

// http://documentcloud.github.com/backbone/#Model-parse
parse : function(resp, xhr) {
  return resp;
},

// snip...
});
```

If you need to accept JSON with a root node, you can override `parse` in each of your models, or override the prototype's function. You'll need to override it on the appropriate collection(s), too.

If you need to send JSON back to the server that includes a root node, you can override `toJSON`, per-model or across all models. When you do this, you'll need to explicitly specify the name of the root key. We use a convention of a `modelName` function on your model to provide this:

```
Backbone.Model.prototype.toJSON = function() {
  var hashWithRoot = {};
  hashWithRoot[this.modelName] = this.attributes;
  return _.clone(hashWithRoot);
};

var Task = Backbone.Model.extend({
  modelName: "task",

  // ...
});
```

5.6. Converting an existing page/view area to use Backbone.js

We'll cover Backbone.js Views and Templates in more detail in the Routers, Views, and Templates chapter, but this section is meant to get you started understanding how Backbone.js views work by illustrating the conversion of a Rails view to a Backbone.js view.

It's important to note that a Rails view is not directly analogous to a Backbone.js view. A Rails view is more like a Backbone.js template, and Backbone.js views are more like Rails controllers. This can cause confusion with developers just started with Backbone.js.

Consider the following Rails view for a tasks index.

```
<h1>Tasks</h1>

<table>
  <tr>
    <th>Title</th>
    <th>Completed</th>
  </tr>

  <% @tasks.each do |task| %>
    <tr>
```

```
<td><%= task.title %></td>
<td><%= task.completed %></td>
</tr>
<% end %>
</table>
```

Assuming we have the Backbone.js Task model and collection and the Rails Task model and controller discussed above, and we're priming the pump with all the tasks, before we can convert the template we must create a Backbone.js view which will render the Backbone.js template.

A Backbone.js view is a class that is responsible for rendering the display of a logical element on the page. A view can also bind to events which may cause it to be re-rendered. For more detailed coverage of Backbone.js views, see the Routers, Views, and Templates chapter.

The most rudimentary view we could introduce at this point would be one that merely renders the above page markup, looping over each task in the Tasks collection. While this would be insufficient for most actual applications, in order to illustrate the building blocks of a Backbone.js view, such a view would be like the one shown below.

```
ExampleApp.Views.TasksIndex = Backbone.View.extend({
  initialize: function() {
    this.render();
  },

  render: function () {
    $(this.el).html(JST['tasks/index']({ tasks: ExampleApp.tasks }));
    $('body').html(this.el);

    return this;
  }
});
```

The Backbone.js view above has an initialize method which will be called when the view is instantiated. This initialize method calls the render method of the view. It's not necessary to immediately render upon initialization, but it's fairly common to do so.

The render method above then renders the *tasks/index* template, passing the collection of tasks into the template. It then sets the HTML of the body element of the page to be the rendered template.

Each Backbone.js view has an element which is stored in this.el. This element can be populated with content, but isn't on the page until placed there by you.

Finally, the Router must be changed to instantiate this view, as shown in the following Tasks router.

```
ExampleApp.Routers.Tasks = Backbone.Router.extend({
  routes: {
    "": "index"
  },

  index: function() {
    new ExampleApp.Views.TasksIndex();
  }
});
```

Now that we have the Backbone.js view in place that renders the template, and its being called by the router, we can focus on converting the above Rails view to a Backbone.js template.

Backbone.js depends on Underscore.js which provides templating. Fortunately, the delimiter and basic concepts used for both Underscore.js and Erb are the same, making conversion relatively painless. For this reason, we recommend using Underscore.js templates when converting a larger, existing Rails application to Backbone.js.

The tasks index template does two things:

- Loops over all of the tasks
- For each task, it outputs the task title and completed attributes

Underscore.js provides many iteration functions that will be familiar to Rails developers. For example, each, map, and reject. Fortunately, Backbone.js also proxies to Underscore.js to provide 26 iteration functions on Backbone.Collection. This means that its possible to call the Underscore.js methods directly on Backbone.js collections.

So we'll use the each method to iterate through the Tasks collection that was passed to the view, as shown in the converted Rails template, which is now an Underscore.js template, below.

```
<h1>Tasks</h1>

<table>
  <tr>
    <th>Title</th>
    <th>Completed</th>
  </tr>

  <% tasks.each(function(model) { %>
    <tr>
      <td><%= model.escape('title') %></td>
      <td><%= model.escape('completed') %></td>
    </tr>
  <% }); %>
</table>
```

As you can see above in the above example, the same delimiter, and the use of the each method make the conversion of the Rails view to an Underscore.js template straightforward.

Finally, in Rails 3.0 and above template output is escaped. In order to ensure that we have the same XSS protection as we did in our Rails template, we access and output the Backbone.js model attributes using the escape method instead of the normal get method.

5.6.1. Breaking out the TaskView

As mentioned above, this simple conversion of the index which merely loops over each of the tasks is not one you'd likely see in a real Backbone.js application.

Backbone.js views should represent the logic pieces of your web page. In the above example, we have an index view, which is a logic piece, but then it is made up of the display of individual tasks. Each of those individual tasks should be represented by a new Backbone.js view, named TaskView.

The benefit of this logical separation is covered in more detail in the Views section, but know that one of the major features of Backbone.js is event binding. With each of the Task models represented by an

individual task view, when that individual model changes the view can be re-rendered automatically (by triggering events) and the entire page doesn't need to be re-rendered.

Continuing our task index example from above, a TaskView will be responsible for rendering just the individual table row for a Task, therefore, its template will appear as follows.

```
<tr>
  <td><%= model.escape('title') %></td>
  <td><%= model.escape('completed') %></td>
</tr>
```

And the Task index template will be changed to be as shown below.

```
<h1>Tasks</h1>

<table>
  <tr>
    <th>Title</th>
    <th>Completed</th>
  </tr>

</table>
```

As you can see above in the index template, the individual tasks are no longer iterated over and rendered inside the table. This will now happen in the TasksIndex and TaskView view, which is shown below.

```
ExampleApp.Views.TaskView = Backbone.View.extend({
  initialize: function() {
  },

  render: function () {
    $(this.el).html(JST['tasks/view']({ model: this.model }));
    return this;
  }
});
```

The TaskView view above is very similar to the one we saw previously for the TasksIndex view. However, unlike the TasksIndex view, the TaskView does not insert itself into the dom. Instead, it only inserts its content into it's own element and the TasksIndex view be responsible for inserting the rendered task into the dom, as shown below.

```
ExampleApp.Views.TasksIndex = Backbone.View.extend({
  initialize: function() {
    this.render();
  },

  render: function () {
    $(this.el).html(JST['tasks/index']({ tasks: ExampleApp.tasks }));

    var tasksIndexView = this;
    ExampleApp.tasks.each(function(task) {
      var taskView = new ExampleApp.Views.TaskView({model: task});
      tasksIndexView.$('table').append(taskView.render().el);
    });

    $('body').html(this.el);
  }
});
```

```
    return this;
  }
});
```

In the new `TasksIndex` view above, the `Tasks` collection is iterated over. For each task, a new `TaskView` is instantiated, rendered, and then inserted into the dom.

If you take a look at the output of the `TasksIndex`, it will appear as follows.

```
<div>
  <h1>Tasks</h1>

  <table>
    <tr>
      <th>Title</th>
      <th>Completed</th>
    </tr>

    <div>
      <tr>
        <td>Task 1</td>
        <td>true</td>
      </tr>
    </div>
    <div>
      <tr>
        <td>Task 2</td>
        <td>false</td>
      </tr>
    </div>
  </table>
</div>
```

Unfortunately, we can see that there is a problem with the above rendered view, and that is the surrounding `div` around each of the rendered tasks.

Each of the rendered tasks has a surrounding `div` because this is the element that each view has that is accessed via `this.el`, and what the view's content is inserted into. By default, this element is a `div` and therefore every view will be wrapped in an extra `div`. While sometimes this extra `div` doesn't really matter, as in the outermost `div` that wraps the entire index, other times this produced invalid markup.

Fortunately, `Backbone.js` provides us with a clean and simple mechanism for changing the element to something other than a `div`. In the case of the `TaskView`, we would like this element to be a `tr`, then the wrapping `tr` can be removed from the task view template.

The element to use is specified by the `tagName` member of the `TaskView`, as shown below.

```
ExampleApp.Views.TaskView = Backbone.View.extend({
  tagName: "tr",

  initialize: function() {
  },

  render: function () {
    $(this.el).html(JST['tasks/view']({ model: this.model }));
    return this;
  }
});
```

Given the above tagName customization, the task view template will be as follows.

```
<td><%= model.escape('title') %></td>
<td><%= model.escape('completed') %></td>
```

And the resulting output of the TasksIndex will be much cleaner, as shown below.

```
<div>
  <h1>Tasks</h1>

  <table>
    <tr>
      <th>Title</th>
      <th>Completed</th>
    </tr>

    <tr>
      <td>Task 1</td>
      <td>true</td>
    </tr>
    <tr>
      <td>Task 2</td>
      <td>false</td>
    </tr>
  </table>
</div>
```

That is the basic building blocks of converting Rails views to Backbone.js and getting a functional system. The majority of Backbone.js programming you will do will likely be in the Views and Templates and there is a lot more too them: event binding, different templating strategies, helpers, event unbinding, and more. All of which are covered in the Routers, Views, and Templates chapter.

5.7. Automatically using the Rails authentication token

When using Backbone.js in a Rails app, you will run into a conflict with the Rails built in Cross Site Scripting (XSS) protection.

When Rails XSS is enabled, each POST or PUT request to Rails should include a special token which is verified to ensure that the request originated from a user which is actually using the Rails app. In recent versions of Rails, Backbone.js Ajax requests are no exception.

To get around this, you have two options. Disable Rails XSS protection (not recommended), or make Backbone.js play nicely with Rails XSS.

To make Backbone.js play nicely with Rails XSS you can monkeypatch Backbone.js to include the Rails XSS token in any requests it makes.

The following is one such script.

```
//
// With additions by Maciej Adwent http://github.com/Maciek416
// If token name and value are not supplied, this code Requires jQuery
//
// Adapted from:
```

```
// http://www.ngauthier.com/2011/02/backbone-and-rails-forgery-protection.html
// Nick Gauthier @ngauthier
//

var BackboneRailsAuthTokenAdapter = {

  //
  // Given an instance of Backbone, route its sync() function so that
  // it executes through this one first, which mixes in the CSRF
  // authenticity token that Rails 3 needs to protect requests from
  // forgery. Optionally, the token's name and value can be supplied
  // by the caller.
  //
  fixSync: function(Backbone, paramName /*optional*/, paramValue /*optional*/){

    if(typeof(paramName)=='string' && typeof(paramValue)=='string'){
      // Use paramName and paramValue as supplied
    } else {
      // Assume we've rendered meta tags with erb
      paramName = $("meta[name='csrf-param']").attr('content');
      paramValue = $("meta[name='csrf-token']").attr('content');
    }

    // alias away the sync method
    Backbone._sync = Backbone.sync;

    // define a new sync method
    Backbone.sync = function(method, model, success, error) {

      // only need a token for non-get requests
      if (method == 'create' || method == 'update' || method == 'delete') {

        // grab the token from the meta tag rails embeds
        var auth_options = {};
        auth_options[paramName] = paramValue;

        // set it as a model attribute without triggering events
        model.set(auth_options, {silent: true});
      }

      // proxy the call to the old sync method
      return Backbone._sync(method, model, success, error);
    };
  },

  // change Backbone's sync function back to the original one
  restoreSync: function(Backbone){
    Backbone.sync = Backbone._sync;
  }
};

BackboneRailsAuthTokenAdapter.fixSync(Backbone);
```

The above patch depends on jQuery, and should be included in your after jQuery and Backbone.js are loaded. Using Jammit, you'd list it below the backbone.js file.

In Rails 3.1, you'll place this file in lib/assets/javascripts. In the example app, you can find this in `example_app/lib/assets/javascripts/backbone.authtokenadapter.js`.

6. Routers, Views, and Templates

6.1. View explanation (chapter unstated)

6.2. Routers (chapter unstated)

6.3. Templating strategy (chapter unstated)

6.4. View helpers (chapter unstated)

6.5. Form helpers (chapter unstated)

6.6. Event binding

A big part of writing snappy rich client applications is building models and views that update in realtime with respect to one another. With Backbone.js you accomplish this with events.

TODO: This is probably the first time we dive into events, unless we touch on them earlier in the models/collections sections. Might want to introduce the topic with a basic example that uses `Backbone.Events` without views & models.

There are three primary kinds of events that your views will bind to:

- DOM events within the view's `this.el` element
- Backbone events triggered by the view's model or collection
- Custom view events

TODO: This three-point breakdown is the wrong way to slice this. Instead of "DOM, model/collection, custom" it should be "DOM, events I observe, events I publish". Events that your view observes need to be cleaned up upon disposing the view, regardless of where those events are triggered (models, collections, or other views, or other arbitrary objects). Events that your view publishes need to be handled in a different way.

TODO: Consider promoting events and binding/unbinding to its own top-level section; this isn't view-specific, although the view layer is where you'll be doing most of your binding.

6.6.1. Binding to DOM events within the view element

The primary function of a view class is to provide behavior for its markup's DOM elements. You can attach event listeners by hand if you like:

```
<!-- templates/soundboard.jst -->
<a class="sound">Honk</a>
```



```
<a class="sound">Beep</a>
```

```
var SoundBoard = Backbone.View.extend({
  render: function() {
    $(this.el).html(JST['soundboard']());
    this.$("a.sound").bind("click", this.playSound);
  },

  playSound: function() {
    // play sound for this element
  }
});
```

But Backbone provides an easier and more declarative approach with the `events` hash:

```
var SoundBoard = Backbone.View.extend({
  events: {
    "click a.sound": "playSound"
  },

  render: function() {
    $(this.el).html(JST['soundboard']());
  },

  playSound: function() {
    // play sound for this element
  }
});
```

Backbone will bind the events with the `Backbone.View.prototype.delegateEvents()` [<http://documentcloud.github.com/backbone/#View-delegateEvents>] function. It binds DOM events with `$.delegate()`, whether you're using the jQuery [<http://api.jquery.com/delegate/>] or Zepto [<https://github.com/madrobby/zepto/blob/v0.7/src/event.js#L96-108>] `.delegate()` function.

It also takes care of binding the event handlers' `this` to the view instance using `_.bind()`.

6.6.2. Binding to events triggered by `this.model` or `this.collection`

In almost every view you write, the view will be bound to a `Backbone.Model` or `Backbone.Collection`, most often with the convenience properties `this.model` or `this.collection`.

TODO: Make sure we discussed the convenience properties previously?

Consider a view that displays a collection of `Task` models. It will re-render itself when any model in the collection is changed or removed, or when a new model is added:

```
var TasksIndex = Backbone.View.extend({
  template: JST['tasks/tasks_index'],
  tagName: 'section',
  id: 'tasks',

  initialize: function() {
    _.bindAll(this, "render");
    this.collection.bind("change", this.render);
    this.collection.bind("add", this.render);
    this.collection.bind("remove", this.render);
  }
});
```

```
    },  
  
    render: function() {  
      $(this.el).html(this.template({tasks: this.collection}));  
    }  
  });  
});
```

6.6.3. Binding to custom events

With sufficiently complex views, you may encounter a situation where you want one view to change in response to another.

TODO: Expound on this situation, discuss that it's unlikely, and you should consider whether you should be binding to models instead. However, sometimes it's useful.

Consider a simple example with a table of users and a toggle control that filters the users to a particular gender:

```
GenderFilter = Backbone.View.extend({  
  events: {  
    "click .show-male": "showMale",  
    "click .show-female": "showFemale",  
    "click .show-both": "showBoth"  
  },  
  
  showMale: function() { this.trigger("changed", "male"); },  
  showFemale: function() { this.trigger("changed", "female"); },  
  showBoth: function() { this.trigger("changed", "both"); }  
});  
  
UsersTable = Backbone.View.extend({  
  initialize: function() {  
    this.filterView = new UserFilter();  
    this.filterView.bind("changed", this.filterByGender);  
  },  
  
  filterByGender: function(gender) {  
    this.filteredCollection = this.collection.byGender(gender);  
  }  
});
```


GenderFilter is responsible for the filter control, and triggers an event with `Backbone.Events.prototype.trigger()` when it changes. UsersTable observes this event, and filters its own collection in response.

6.7. Cleaning Up: Unbinding


In the last section, we discussed three different kinds of event binding in your `Backbone.Views` classes: DOM events, model/collection events, and custom view events. Next we'll discuss unbinding these events: why it's a good idea, and how to do it.

6.7.1. Why do I have to unbind events?

Consider two views in a Todo app: an index view which contains all the tasks that need to be done:

Figure 2. Tasks index view

and a detail view that shows detail on one task:

Figure 3. Tasks detail view

The interface switches between the two views.

Here's the source for the aggregate index view:

```
var TasksIndex = Backbone.View.extend({
  template: JST['tasks/tasks_index'],
  tagName: 'section',
  id: 'tasks',

  initialize: function() {
    _.bindAll(this, "render");
    this.collection.bind("change", this.render);
    this.collection.bind("add", this.render);
    this.collection.bind("remove", this.render);
  },

  render: function() {
    $(this.el).html(this.template({tasks: this.collection}));
  }
});
```

and the source for the individual task detail view:

```
var TaskDetail = Backbone.View.extend({
  template: JST['tasks/tasks_detail'],
  tagName: 'section',
  id: 'task',

  events: {
    "click .comments .form-inputs button": "createComment"
  },

  initialize: function() {
    _.bindAll(this, "render");

    this.model.bind("change", this.render);
    this.model.comments.bind("change", this.render);
    this.model.comments.bind("add", this.render);
  },

  render: function() {
    $(this.el).html(this.template({task: this.model}));
  },

  createComment: function() {
    var comment = new Comment({ text: this.$('.new-comment-input').val() });
    this.$('.new-comment-input').val('');
    this.model.comments.create(comment);
  }
});
```

Each task on the index page links to the detail view for itself. When a user follows one of these links and navigates from the index page to the detail page, then interacts with the detail view to change a model, the `change` event on the `TaskApp.tasks` collection is fired. One consequence of this is that the index view, which is still bound and observing the `change` event, will re-render itself.

This is both a functional bug and a memory leak: not only will the index view re-render and disrupt the detail display momentarily, but navigating back and forth between the views without disposing of the previous view will keep creating more views and binding more events on the associated models or collections.

These can be extremely tricky to track down on a production application, especially if you are nesting child views. Sadly, there's no "garbage collection" for views in Backbone, so your application needs to manage this itself.

Let's take a look at how to unbind various kinds of events.

6.7.2. Unbinding DOM events

When you call `this.remove()` in your view, it delegates to `jQuery.remove()` by invoking `$(this.el).remove()`. This means that jQuery takes care of cleaning up any events bound on DOM elements within your view, regardless of whether you bound them with the Backbone `events` hash or by hand; for example, with `$.bind()`, `$.delegate()`, or `$.live()`.

6.7.3. Unbinding model and collection events

If your view binds to events on a model or collection, you are responsible for unbinding these events. You do this with a simple call to `this.model.unbind()` or `this.collection.unbind()`; the

`Backbone.Events.unbind()` function [<http://documentcloud.github.com/backbone/#Events-unbind>] removes all callbacks on that object.

When should we unbind these handlers? Whenever the view is going away. This means that any pieces of code that create new instances of this view become responsible for cleaning up after it. That doesn't sound like a very cohesive approach, so let's include the cleanup responsibility on this view.

TODO: Consider just overriding `Backbone.View.prototype.remove()` instead of making a new function, since `remove()` is very simple. What are the pros/cons?

Let's write a `leave()` function on our view that wraps `remove()` and handles any additional event unbinding we need to do. As a convention, when we use this view elsewhere, we'll call `leave()` instead of `remove()` when we're done:

```
var SomeCollectionView = Backbone.View.extend({
  // snip...

  initialize: function() {
    this.collection.bind("change", this.render);
  },

  leave: function() {
    this.collection.unbind("change", this.render);
    this.remove();
  }

  // snip...
});
```

6.7.4. Keep track of `bind()` calls to unbind more easily

In the example above, unbinding the collection change event isn't too much hassle; since we're only observing one thing, we only have to unbind one thing. But even the addition of one line to `leave()` is easy to forget, and if you bind to multiple events then it only gets more verbose.

Let's add a step of indirection in event binding so that we can automatically clean up all the events with one call. We'll add and use a `bindTo()` function that keeps track of all the event handlers we bind, and then issue a single call to `unbindFromAll()` to unbind them:

```
var SomeCollectionView = Backbone.View.extend({
  initialize: function() {
    this.bindTo(this.collection, "change", this.render);
  },

  leave: function() {
    this.unbindFromAll();
    this.remove();
  },

  bindTo: function(source, event, callback) {
    source.bind(event, callback, this);
    this.bindings.push({ source: source, event: event, callback: callback });
  },

  unbindFromAll: function() {
    _.each(this.bindings, function(binding) {
      binding.source.unbind(binding.event, binding.callback);
    });
  }
});
```

```
});  
  this.bindings = [];  
}  
});
```

These functions, `bindTo()` and `unbindFromAll()`, can be extracted into a reusable mixin or superclass. Then, we just have to use `bindTo()` instead of `model.bind()` and be assured that the handlers will be cleaned up during `leave()`.

TODO: Is it viable to use `Function.caller` inside `Backbone.Events` so this functionality is provided by `Backbone.Events`? <https://gist.github.com/158a4172aea28876d0fc>

TODO: Wrap `bindTo()` and `unbindFromAll()` into `Observer` which gets mixed into `CompositeView`.

6.7.5. Unbinding custom events

With the first two kinds of event binding that we discussed, DOM and model/collection, the view is the observer. The responsibility to clean up is on the observer, and here the responsibility consists of unbinding the event handler when the view is being removed.

But other times, our view classes will trigger (emit) events of their own. In this case, other objects are the observer, and are responsible for cleaning up the event binding when they are disposed.

However, additionally, when the view itself is disposed of with `leave()`, it should clean up any event handlers bound on **itself** for events that it triggers.

This is handled by invoking `Backbone.Events.unbind()`:

```
var FilteringView = Backbone.View.extend({  
  // snip...  
  
  events: {  
    "click a.filter": "changeFilter"  
  },  
  
  changeFilter: function() {  
    if (someLogic()) {  
      this.trigger("filtered", { some: options });  
    }  
  },  
  
  leave: function() {  
    this.unbind(); // Clean up any event handlers bound on this view  
    this.remove();  
  }  
  
  // snip...  
});
```

6.7.6. Establish a convention for consistent and correct unbinding

There's no built-in garbage collection for Backbone's event bindings, and forgetting to unbind can cause bugs and memory leaks. The solution is to make sure you unbind events and remove views

when you leave them. Our approach to this is two-fold: write a set of reusable functions that manage cleaning up a view's bindings, and use these functions where ever views are instantiated: in `Router` instances, and in composite views. We'll take a look at these concrete, reusable approaches in the next two sections about `SwappingRouter` and `CompositeView`.

6.8. Swapping router

When switching from one view to another, we should clean up the previous view. We discussed previously a convention of writing a `view.leave()` Let's augment our view to include the ability to clean itself up by "leaving" the DOM:

```
var MyView = Backbone.View.extend({
  // ...

  leave: function() {
    this.unbind();
    this.remove();
  },

  // ...
});
```

The `unbind()` and `remove()` functions are provided by `Backbone.Events` and `Backbone.Events.unbind()` will remove all callbacks registered on the view, and `remove()` will remove the view's element from the DOM.

In simple cases, we replace one full page view with another full page (less any shared layout). We introduce a convention that all actions underneath one `Router` share the same root element, and define it as `el` on the router.

Now, a `SwappingRouter` can take advantage of the `leave()` function, and clean up any existing views before swapping to a new one. It swaps into a new view by rendering that view into its own `el`:

```
SwappingRouter = function(options) {
  Backbone.Router.apply(this, [options]);
};

_.extend(SwappingRouter.prototype, Backbone.Router.prototype, {
  swap: function(newView) {
    if (this.currentView && this.currentView.leave) {
      this.currentView.leave();
    }

    this.currentView = newView;
    this.currentView.render();
    $(this.el).empty().append(this.currentView.el);
  }
});

SwappingRouter.extend = Backbone.Router.extend;
```

Now all you need to do in a route function is call `swap()`, passing in the new view that should be rendered. The `swap()` function's job is to call `leave()` on the current view, render the new view and append it to the router's `el`, and finally store who the current view is, so that next time `swap()` is invoked, it can be properly cleaned up as well.

6.8.1. SwappingRouter and Backbone internals

If the code for `SwappingRouter` seems a little confusing, don't fret: it is, thanks to JavaScript's object model! Sadly, it's not as simple to just drop in the `swap` method into `Backbone.Router`, or call `Backbone.Router.extend` to mixin the function we need.

Our goal here is essentially to create a subclass of `Backbone.Router`, and to extend it without modifying the original class. This gives us a few benefits: first, `SwappingRouter` should work with Backbone upgrades. Second, it should be **obvious** and **intention-revealing** when a controller needs to swap views. If we chose to just mix in a `swap` method, and called it from a direct descendant of `Backbone.Router`, an unaware (and unlucky) programmer now needs to go on a deep source dive in an attempt to figure out where that's coming from. At least with a subclass, the hunt should start at the file where it was defined.

The procedure used to create `SwappingRouter` is onerous thanks to a mix of Backbone-isms and just how clunky inheritance is in JavaScript. First off, we need to define the constructor, which delegates to the `Backbone.Router` constructor with the use of `Function#apply`. The next block of code uses Underscore's `Object#extend` to create the set of functions and properties that will become `SwappingRouter`. The `extend` function takes a destination, in this case the empty prototype for `SwappingRouter`, and copies in the properties in the `Backbone.Router` prototype along with our new custom object that includes the `swap` function.

Finally, the subclass cake is topped off with some Backbone frosting: setting `extend`, which is a self-propagating function that all Backbone public classes use. Let's take a quick look at this function, as of Backbone 0.5.3:


```
var extend = function (protoProps, classProps) {
  var child = inherits(this, protoProps, classProps);
  child.extend = this.extend;
  return child;
};

// Helper function to correctly set up the prototype chain, for subclasses.
// Similar to `goog.inherits`, but uses a hash of prototype properties and
// class properties to be extended.
var inherits = function(parent, protoProps, staticProps) {
  // sparing our readers the internals of this function... for a deep dive
  // into the dark realms of JavaScript's prototype system, read the source!
}
```

So, it's a function that calls `inherits` to make a new subclass. The comments reference `goog.inherits` from Google's Closure Library, which contains similar utility functions to allow more class-style inheritance.

The end result here is that whenever you make a custom controller, internally in Backbone, you're making **another** subclass. The inheritance chain for `TasksRouter` would then look like:

Figure 4. Router class inheritance



Phew! Hopefully this adventure into Backbone and JavaScript internals has taught you that although it's more code, it's hopefully going to save time down the road for those maintaining your code.

6.9. Composite views

The `SwappingRouter` above calls `leave()` on the view it currently holds. This function is not part of Backbone itself, and is part of our extension library to help make views more modular and maintainable. This section goes over the Composite View pattern, the `CompositeView` class itself, and some concerns to keep in mind while creating your views.

6.9.1. Refactoring from a large view

One of the first refactorings you find yourself doing in a non-trivial Backbone app is splitting up large views into composable parts. Let's take another look at the `TaskDetail` source code from the beginning of this section:

```
var TaskDetail = Backbone.View.extend({
  template: JST['tasks/tasks_detail'],
  tagName: 'section',
  id: 'task',

  events: {
    "click .comments .form-inputs button": "createComment"
  },

  initialize: function() {
    _.bindAll(this, "render");

    this.model.bind("change", this.render);
    this.model.comments.bind("change", this.render);
    this.model.comments.bind("add", this.render);
  },

  render: function() {
    $(this.el).html(this.template({task: this.model}));
  },

  createComment: function() {
    var comment = new Comment({ text: this.$('.new-comment-input').val() });
    this.$('.new-comment-input').val('');
    this.model.comments.create(comment);
  }
});
```

The view class references a template, which renders out the HTML for this page:

```
<section class="task-details">
  <input type="checkbox"<%= task.isComplete() ? ' checked="checked"' : '' %> />
  <h2><%= task.escape("title") %></h2>
</section>

<section class="comments">
  <ul>
    <% task.comments.each(function(comment) { %>
      <li>
        <h4><%= comment.user.escape('name') %></h4>
        <p><%= comment.escape('text') %></p>
```

```

    </li>
    <% } %>
  </ul>

  <div class="form-inputs">
    <label for="new-comment-input">Add comment</label>
    <textarea id="new-comment-input" cols="30" rows="10"></textarea>
    <button>Add Comment</button>
  </div>
</section>

```

There are clearly several concerns going on here: rendering the task, rendering the comments that folks have left, and rendering the form to create new comments. Let's separate those concerns. A first approach might be to just break up the template files:

```

<!-- tasks/show.jst -->
<section class="task-details">
  <%= JST['tasks/details']({ task: task }) %>
</section>

```

```

<section class="comments">
  <%= JST['comments/list']({ task: task }) %>
</section>

```

```

<!-- tasks/details.jst -->
<input type="checkbox"<%= task.isComplete() ? ' checked="checked"' : '' %> />
<h2><%= task.escape("title") %></h2>

```

```

<!-- comments/list.jst -->
<ul>
  <% task.comments.each(function(comment) { %>
    <%= JST['comments/item']({ comment: comment }) %>
  <% } %>
</ul>

```

```

<%= JST['comments/new']() %>

```

```

<!-- comments/item.jst -->
<h4><%= comment.user.escape('name') %></h4>
<p><%= comment.escape('text') %></p>

```

```

<!-- comments/new.jst -->
<div class="form-inputs">
  <label for="new-comment-input">Add comment</label>
  <textarea id="new-comment-input" cols="30" rows="10"></textarea>
  <button>Add Comment</button>
</div>

```

But this is really only half the story. The `TaskDetail` view class still handles multiple concerns: displaying the task, and creating comments. Let's split that view class up, using the `CompositeView` base class:

```

CompositeView = function(options) {
  this.children = [];
  Backbone.View.apply(this, [options]);
};

_.extend(CompositeView.prototype, Backbone.View.prototype, {
  leave: function() {
    this._leaveChildren();
  }
});

```

```

    this.unbind();
    this.remove();
    this._removeFromParent();
  },

  removeChild: function(view) {
    var index = this.children.indexOf(view);
    this.children.splice(index, 1);
  },

  renderChild: function(view) {
    view.render();
    this.children.push(view);
    view.parent = this;
  },

  appendChild: function(view) {
    this.renderChild(view);
    $(this.el).append(view.el);
  },

  renderChildInto: function(view, container) {
    this.renderChild(view);
    $(container).html('').append(view.el);
  },

  _leaveChildren: function() {
    var clonedChildren = this.children.slice(0);
    _each(clonedChildren, function(view) {
      if (view.leave) {
        view.leave();
      }
    });
  },

  _removeFromParent: function() {
    if (this.parent) {
      this.parent.removeChild(this);
    }
  }
});

CompositeView.extend = Backbone.View.extend;

```

TODO: Re-link to swapping-internals anchor once <https://github.com/schacon/git-scribe/issues/33> is fixed

Similar to the `SwappingRouter`, the `CompositeView` base class solves common housekeeping problems by establishing a convention. See the `Swapping Router and Backbone internals` section for an in-depth analysis of how this subclassing pattern works.

Now our `CompositeView` maintains an array of its immediate children as `this.children`. With this reference in place, a parent view's `leave()` method can invoke `leave()` on its children, ensuring that an entire tree of composed views is cleaned up properly.

For child views that can dismiss themselves, such as dialog boxes, children maintain a back-reference at `this.parent`. This is used to reach up and call `this.parent.removeChild(this)` for these self-dismissing views.

Making use of `CompositeView`, we split up the `TaskDetail` view class:

```
var TaskDetail = Backbone.View.extend({
  tagName: 'section',
  id: 'task',

  initialize: function() {
    _.bindAll(this, "renderDetails");
    this.model.bind("change", this.renderDetails);
  },

  render: function() {
    this.renderLayout();
    this.renderDetail();
    this.renderCommentsList();
  },

  renderLayout: function() {
    $(this.el).html(JST['tasks/show']());
  },

  renderDetails: function() {
    var detailsMarkup = JST['tasks/details']({ task: this.model });
    this.$('.task-details').html(detailsMarkup);
  },

  renderCommentsList: function() {
    var commentsList = new CommentsList({ model: this.model });
    var commentsContainer = this.$('comments');
    this.renderChildInto(commentsList, commentsContainer);
  }
});

var CommentsList = CompositeView.extend({
  tagName: 'ul',

  initialize: function() {
    this.model.comments.bind("add", this.renderComments);
  },

  render: function() {
    this.renderLayout();
    this.renderComments();
    this.renderCommentForm();
  },

  renderLayout: function() {
    $(this.el).html(JST['comments/list']());
  },

  renderComments: function() {
    var commentsContainer = this.$('comments-list');
    commentsContainer.html('');

    this.model.comments.each(function(comment) {
      var commentMarkup = JST['comments/item']({ comment: comment });
      commentsContainer.append(commentMarkup);
    });
  },
});
```

```

renderCommentForm: function() {
  var commentForm = new CommentForm({ model: this.model });
  var commentFormContainer = this.$('.new-comment-form');
  this.renderChildInto(commentForm, commentFormContainer);
}
});

var CommentForm = CompositeView.extend({
  events: {
    "click button": "createComment"
  },

  initialize: function() {
    this.model = this.options.model;
  },

  render: function() {
    $(this.el).html(JST['comments/new']);
  },

  createComment: function() {
    var comment = new Comment({ text: $('.new-comment-input').val() });
    this.$('.new-comment-input').val('');
    this.model.comments.create(comment);
  }
});

```

Along with this, remove the `<%= JST(...) %>` template nestings, allowing the view classes to assemble the templates instead. In this case, each template contains placeholder elements that are used to wrap child views:

```

<!-- tasks/show.jst -->
<section class="task-details">
</section>

<section class="comments">
</section>

```

```

<!-- tasks/details.jst -->
<input type="checkbox"<%= task.isComplete() ? ' checked="checked"' : '' %> />
<h2><%= task.escape("title") %></h2>

```

```

<!-- comments/list.jst -->
<ul class="comments-list">
</ul>

<section class="new-comment-form">
</section>

```

```

<!-- comments/item.jst -->
<h4><%= comment.user.escape('name') %></h4>
<p><%= comment.escape('text') %></p>

```

```

<!-- comments/new.jst -->
<label for="new-comment-input">Add comment</label>
<textarea class="new-comment-input" cols="30" rows="10"></textarea>
<button>Add Comment</button>

```

There are several advantages to this approach:

- Each view class has a smaller and more cohesive set of responsibilities.
- The comments view code, extracted and decoupled from the task view code, can now be reused on other domain objects with comments.
- The task view performs better, since adding new comments or updating the task details will only re-render the pertinent section, instead of re-rendering the entire task + comments composite.

6.9.2. Cleaning up views properly

We now have a full set of tools to clean up views properly.

TODO: Wrap up and re-state the "cleaning up, swappingrouter, compositeview" sections. Mix `Observer` into `CompositeView`.

6.10. How to use multiple views on the same model/collection (chapter unstated)

6.11. Internationalization (chapter unstated)

7. Models and collections

7.1. Naming conventions (chapter unstated)

7.2. Nested resources (chapter unstated)

7.3. Model associations

Backbone.js doesn't prescribe a way to define associations between models, so we need to get creative and use the power of JavaScript to set up associations in such a way that its usage is natural.

7.3.1. Belongs to associations

Setting up a `belongs_to` association in Backbone is a two step process. Let's discuss setting up the association that may occur between a task and a user. The end result of the approach is a `Task` instance having a property called `user` where we store the associated `User` object.

To set this up, let's start by telling Rails to augment the task's JSON representation to also send over the associated user attributes:

```
class Task < ActiveRecord::Base
  belongs_to :user

  def as_json(options = {})
    super(include: { user: { only: [:name, :email] } })
  end
end
```

```
end
```

This means that when Backbone calls `fetch()` for a `Task` model, it will include the name and email of the associated user nested within the task JSON representation. Something like this:

```
{
  "title": "Buy more Cheeseburgers",
  "due_date": "2011-03-04",
  "user": {
    "name": "Robert McGraffalon",
    "email": "bobby@themcgraffalons.com"
  }
}
```

Now that we receive user data with the task's JSON representation, let's tell our Backbone User model to store the User object. We do that on the task's initializer. Here's a first cut at that:

```
var Task = Backbone.Model.extend({
  initialize: function() {
    this.user = new User(this.get('user'));
  }
});
```

We can make a couple of improvements to the above. First, you'll soon realize that you might be setting the user outside of the initialize as well. Second, the initializer should check whether there is user data in the first place. To address the first concern, let's create a setter for the object. Backbone provides a handy function called `has` that returns true or false depending on whether the provided option is set for the object:

```
var Task = Backbone.Model.extend({
  initialize: function() {
    if (this.has('user')) {
      this.setUser(new User(this.get('user')));
    }
  },

  setUser: function(user) {
    this.user = user;
  }
});
```

The final setup allows for a nice clean interface to a task's user, by accessing the task property of the user instance.

```
var task = Task.fetch(1);
console.log(task.get('title') + ' is being worked on by ' + task.user.get('name'));
```

7.3.2. Has many associations

You can take a similar approach to set up a `has_many` association on the client side models. This time, however, the object's property will be a Backbone collection.

Following the example, say we need access to a user's tasks. Let's set up the JSON representation on the Rails side first:

```
class User < ActiveRecord::Base
```



```
has_many :tasks

def as_json(options = {})
  super(include: { tasks: { only: [:body, :due_date] } })
end
end
```

Now, on the Backbone `User` model's initializer, let's call the `setTasks` function:

```
var User = Backbone.Model.extend({
  initialize: function() {
    var tasks = new Tasks.reset(this.get('tasks'));
    this.setTasks(tasks);
  },

  setTasks: function(tasks) {
    this.tasks = tasks;
  }
});
```

Note that we are setting the relation to an instance of the `Tasks` collection.

TODO: Let's expand upon this, as it isn't the most flexible solution. (It is a good start.) We are setting the JSON representation of the Rails models to suit the Backbone.js concerns. Additionally, the `Task#as_json` method at the top is concerned with the `User` JSON representation. It should at least delegate to `User#as_json`. Going further, the JSON presentation for consumption by Backbone.js should be completely extracted into the JSON API endpoint controller action, or even a separate presenter class.

7.4. Filters and sorting

When using our Backbone models and collections, it's often handy to filter the collections by reusable criteria, or sort them by several different criteria.

7.4.1. Filters

To filter a `Backbone.Collection`, like with Rails named scopes, define functions on your collections that filter by your criteria, using the `select` function from `Underscore.js`, and return new instances of the collection class. A first implementation might look like this:

```
var Tasks = Backbone.Collection.extend({
  model: Task,
  url: '/tasks',

  complete: function() {
    var filteredTasks = this.select(function(task) {
      return task.get('completed_at') !== null;
    });
    return new Tasks(filteredTasks);
  }
});
```

Let's refactor this a bit. Ideally, the filter functions will reuse logic already defined in your model class:

```
var Task = Backbone.Model.extend({
  isComplete: function() {
    return this.get('completed_at') !== null;
  }
});

var Tasks = Backbone.Collection.extend({
  model: Task,
  url: '/tasks',

  complete: function() {
    var filteredTasks = this.select(function(task) {
      return task.isComplete();
    });
    return new Tasks(filteredTasks);
  }
});
```

Going further, notice that there are actually two concerns in this function. The first is the notion of filtering the collection, and the other is the specific filtering criteria (`task.isComplete()`).

Let's separate the two concerns here, and extract a `filtered` function:

```
var Task = Backbone.Model.extend({
  isComplete: function() {
    return this.get('completed_at') !== null;
  }
});

var Tasks = Backbone.Collection.extend({
  model: Task,
  url: '/tasks',

  complete: function() {
    return this.filtered(function(task) {
      return task.isComplete();
    });
  },

  filtered: function(criteriaFunction) {
    return new Tasks(this.select(criteriaFunction));
  }
});
```

We can extract this function into a reusable mixin, abstracting the `Tasks` collection class using `this.constructor`:

```
FilterableCollectionMixin = {
  filtered: function(criteriaFunction) {
    return new this.constructor(this.select(criteriaFunction));
  }
};

var Task = Backbone.Model.extend({
  isComplete: function() {
    return this.get('completed_at') !== null;
  }
});
```

```
var Tasks = Backbone.Collection.extend(_.extend({
  model: Task,
  url: '/tasks',

  complete: function() {
    return this.filtered(function(task) {
      return task.isComplete();
    });
  }
}, FilterableCollectionMixin));
```

7.4.2. Sorting

The simplest way to sort a `Backbone.Collection` is to define a comparator function. This functionality is built in:

```
var Tasks = Backbone.Collection.extend({
  model: Task,
  url: '/tasks',

  comparator: function(task) {
    return task.dueDate;
  }
});
```

If you'd like to provide more than one sort order on your collection, you can use an approach similar to the `filtered` function above, and return a new `Backbone.Collection` whose comparator is overridden. Call `sort` to update the ordering on the new collection:

```
var Tasks = Backbone.Collection.extend({
  model: Task,
  url: '/tasks',

  comparator: function(task) {
    return task.dueDate;
  },

  byCreatedAt: function() {
    var sortedCollection = new Tasks(this.models);
    sortedCollection.comparator = function(task) {
      return task.createdAt;
    };
    sortedCollection.sort();
    return sortedCollection;
  }
});
```

Similarly, you can extract the reusable concern to another function:

```
var Tasks = Backbone.Collection.extend({
  model: Task,
  url: '/tasks',

  comparator: function(task) {
    return task.dueDate;
  },
```

```
byCreatedAt: function() {
  return this.sortedBy(function(task) {
    return task.createdAt;
  });
},

byCompletedAt: function() {
  return this.sortedBy(function(task) {
    return task.completedAt;
  });
},

sortedBy: function(comparator) {
  var sortedCollection = new Tasks(this.models);
  sortedCollection.comparator = comparator;
  sortedCollection.sort();
  return sortedCollection;
}
});
```

And then into another reusable mixin:

```
var SortableCollectionMixin = {
  sortedBy: function(comparator) {
    var sortedCollection = new this.constructor(this.models);
    sortedCollection.comparator = comparator;
    sortedCollection.sort();
    return sortedCollection;
  }
};

var Tasks = Backbone.Collection.extend(_.extend({
  model: Task,
  url: '/tasks',

  comparator: function(task) {
    return task.dueDate;
  },

  byCreatedAt: function() {
    return this.sortedBy(function(task) {
      return task.createdAt;
    });
  },

  byCompletedAt: function() {
    return this.sortedBy(function(task) {
      return task.completedAt;
    });
  }
}, SortableCollectionMixin));
```

7.5. Client/Server duplicated business logic (chapter unstated)

7.6. Validations (chapter unstated)

7.7. Synchronizing between clients (chapter unstated)

8. Testing (section unstated)

8.1. Full-stack integration testing

8.2. Isolated unit testing

9. The JavaScript language (section unstated)

9.1. Model attribute types and serialization

9.2. Context binding (JS `this`)

9.3. CoffeeScript with Backbone.js

10. Security (stub)

10.1. XSS with JSON bootstrapping (stub)

Use `json2.js` and:

```
<script type="text/json" id="something">
  <%= something.to_json %>
</script>

<script type="text/javascript">
  (function () {
    var something = JSON.parse($('#something').text());
    someJavaScriptFunction(something);
  })();
</script>
```

10.2. XSS in HTML templates (stub)

TODO: Discuss `Backbone.Model.escape`, `_.escape`, defaulting to escape with `<%=` vs `<%=`, escaping in other templating.

11. Performance (stub)

11.1. Dependency choice

Backbone.js defines a `$` variable that defers to jQuery if present.

If you are only targeting mobile platforms, Backbone will handily fall back to Zepto <http://zeptojs.com> for a more lightweight dependency. Zepto is "a minimalist JavaScript framework for mobile WebKit browsers, with a jQuery-compatible syntax." From `backbone.js`:

```
(function(){  
  
  // Initial Setup  
  // -----  
  
  // Save a reference to the global object.  
  var root = this;  
  
  // For Backbone's purposes, jQuery or Zepto owns the `$` variable.  
  var $ = root.jQuery || root.Zepto;  
  
  // ...  
  
}).call(this);
```